

The information content of a limit order book: The case of an FX market[☆]

Roman Kozhan*, Mark Salmon

Warwick Business School, University of Warwick, Scarman Road, Coventry, CV4 7AL, UK

Available online 4 August 2011

Abstract

In this paper we examine the question of whether knowledge of the information contained in a limit order book helps to provide economic value in a simple trading scheme. Given the greater information content of the order book, over simple price information, it might naturally be expected that the order book would dominate. Using Dollar Sterling tick data, we find that despite the in-sample statistical significance of variables describing the structure of the limit order book in explaining tick-by-tick returns, they do not consistently add significant economic value out-of-sample. We show this using a simple linear model to determine trading activity, as well as a model-free genetic algorithm based on price, order flow, and order book information. We also find that the profitability of all trading rules based on genetic algorithms dropped substantially in 2008 compared to 2003 data.

© 2011 Elsevier B.V. All rights reserved.

JEL classification: D81; F31; C53

Keywords: Profitability; Limit order book; High-frequency data; Algorithmic trading

1. Introduction

One important issue in recent market microstructure research has been whether knowledge of the structure of the limit order book is informative regarding future price movements.

[☆] We are grateful to an anonymous referee and Bruce Lehmann (the editor), whose insightful comments greatly improved the paper. We also thank Martin Evans, Matthijs Fleischer, Thierry Foucault, Lawrence Harris, Rich Lyons, Chris Neely, Carol Osler, Richard Payne, Dagfinn Rime, Lucio Sarno, Nick Webber and Paul Weller for their helpful comments and suggestions. We remain responsible for all errors.

*Corresponding author. Tel.: +44 24 7652 4118.

E-mail addresses: Roman.Kozhan@wbs.ac.uk (R. Kozhan), Mark.Salmon@wbs.ac.uk (M. Salmon).

There is a growing body of theoretical work suggesting that limit orders imply the predictability of short-term asset returns (see [Handa and Schwartz, 1996, 2003](#); [Harris, 1998](#); [Parlour, 1998](#); [Foucault, 1999](#); [Rosu, 2010](#) among others). This is in contrast with earlier papers that implied that informed traders would only use market orders (see [Glosten, 1994](#); [Rock, 1996](#); [Seppi, 1997](#)). This debate has also been carried out empirically by [Harris and Hasbrouck \(1996\)](#), [Kavajecz \(1999\)](#), [Harris and Panchapagesan \(2005\)](#), [Cao, Hansch, and Wang \(2009\)](#), and [Hellström and Simonsen \(2009\)](#), all of whom demonstrated that asset returns can be explained by limit order book information, such as depth and order flow. However, these studies have failed to demonstrate that the predictability of returns can be exploited in economic terms. In this paper we go beyond statistical significance and consider the economic value of limit order book information in an FX market.

We address this question by explicitly constructing trading strategies based on full limit order book and price information in the FX market. These strategies only use historical information in order to ensure that trading can be implemented in “real time” and focus on the economic value of ex ante predictability in out-of-sample prediction exercises.

Such an approach requires the explicit specification of a forecasting model, which might be misspecified and so spuriously unable to exploit the available information. In order to address this issue, we construct trading strategies in a model-free way by employing a genetic algorithm. Genetic algorithms serve as a systematic search mechanism for the best trading rule from amongst a huge universe of potential rules given the particular information set and have been successfully applied in a number of financial applications, most notably by [Dworman, Kimbrough, and Laing \(1996\)](#), [Chen and Yeh \(1997a, 1997b\)](#), [Neely, Weller, and Ditmar \(1997\)](#), [Allen and Karjalainen \(1999\)](#), [Neely and Weller \(2001\)](#), [Dempster and Jones \(2001\)](#), [Chen, Duffy, and Yeh \(1999\)](#), [Arifovic \(1996\)](#). Rather than adopting a single specific forecasting model, the genetic algorithm searches from a very large set for that trading rule which exploits the information most profitably. We then test if this approach generates significantly higher returns when new information constructed from the limit order book is included alongside price information.

It is important to recognize the theoretical and practical coherence offered by using genetic algorithms. A number of authors, since [Leitch and Tanner \(1991\)](#), have argued that the use of purely statistical criteria to evaluate forecasts and trading strategies is inappropriate (e.g., [Satchell and Timmermann, 1995](#); [Granger and Pesaran, 2000](#); [Pesaran and Skouras, 2002](#); [Granger and Machina, 2006](#)). The issue turns on the appropriate loss function and whereas many statistical evaluation criteria are based on a quadratic loss, practical criteria are more likely to be based on the utility derived from profits. Critically, from our point of view, the genetic algorithm constructs trading rules using the same loss function as is used to evaluate the out-of-sample performance of the trading strategy, unlike a linear regression model where a statistical quadratic loss is used in estimation.

Another important factor that needs to be considered when testing the profitability of trading strategies in “real time” is transaction costs.¹ We analyze the performance of our trading rules on the basis of the best bid and ask prices using tick-by-tick data and so explicitly take into account transaction costs as measured by the bid–ask spread. This allows us to test if predictable components in exchange rate returns are economically exploitable net of transaction costs.

¹[Neely and Weller \(2003\)](#) for instance emphasize the critical role of transaction costs and inconsistencies between the data used by practitioners and in academic simulations.

Using data on the U.S. dollar sterling exchange rate for five separate weeks² we find statistical predictability in the exchange rate and profitability net of transaction costs for samples drawn from 2003. However, we find that the profitability in more recent data from 2008 decreases substantially and in most cases is not significantly different from zero. This could be explained by the tremendous recent growth in high-frequency algorithmic trading within financial markets.

We also find in-sample statistical significance of limit order book information in all sample periods. Specifically, we show that both static information about liquidity beyond the best prices and order flow of both market and limit orders have some ability to explain future short-term movements of the exchange rate. However, we find little or no value in an economic sense in allowing the predictor to exploit information in the order book beyond of that contained in the best prices. In other words, we fail to significantly increase out-of-sample returns from our trading strategy when we use liquidity and order flow information. Our main finding then is that any information contained in limit orders beyond best prices is not robust enough to be exploited profitably out-of-sample, particularly in the most recent 2008 data.

The remainder of the paper is organized as follows. In [Section 2](#) we provide a literature review relevant to this research. [Section 3](#) contains a description of the data used in the study and the methodology employed in the analysis. The main results are given in [Section 4](#) and [5](#) concludes.

2. Literature review

Limit order book markets potentially offer greater transparency when compared with quote-driven markets. Whereas dealer markets will usually only release the dealers' best quotes, a limit order book can allow its users to view the depth at a number of price levels away from the market price. The NYSE, under the OpenBook program, publishes aggregate depths at all price levels on either side of the book and under LiquidityQuote displays a bid and offer quote, potentially different from the best quotes in the market. NASDAQ's SuperMontage order entry and execution system displays aggregate depths at five best price levels on either side and employs a scan function that allows traders to assess liquidity further along the book. The question is how this incremental information on the structure of a limit order book is used and whether it adds economic value in the process of price discovery.

There are two closely related literatures that bear on research. The first considers the choice of order type, market or limit order, and then how the structure of the limit order book and traders' preferences affect this choice. The second considers whether the structure of the limit order book is informative regarding the evolution of future prices.

Theoretical papers by [Glosten \(1994\)](#), [Rock \(1996\)](#), and [Seppe \(1997\)](#) assume that informed traders use market orders and so limit orders beyond the best prices can contain little information. Informed traders enter the market to exploit their private information using market orders with guaranteed immediate execution. [Chakravarty and Holden \(1995\)](#) consider a model where informed traders are allowed to submit both limit and market orders and show that an optimal order placement strategy consists of a combination of limit and market orders. [Parlour \(1998\)](#) presents a dynamic model of a limit order market in which the decision to submit a market order or a limit order depends

²We have examined data from a number of different periods and find similar results for all periods.

on the current state of liquidity and the trader's place in the limit order queue. Foucault (1999) presents a dynamic model of order placement in a market with heterogeneous asset valuation among traders with no private information. He showed that there is a trade-off between limit and market order strategies that depends on the volatility of asset returns and bid-ask spreads. Handa and Schwartz (1996, 2003) examine the impact of asymmetric information on order placement strategies. They show that if the cost of being picked-off by an informed trader is lower than the expected gain to limit order execution, then a limit order strategy can be profitable. Ranaldo (2004) examines how the state of the limit order book affects a trader's strategy. He shows that patient traders become more aggressive and hence use more market orders when their side of the book is thicker, the spread wider, and the volatility increases. More recently, Kaniel and Liu (2006) present a simple equilibrium model to investigate whether informed traders use limit or market orders. They show that informed traders prefer limit orders when the private information is long-lived, and hence limit orders convey more information than market orders. When this is the case and the number of traders who can discover the private information is small, then using market orders will reveal too much information implying higher trading costs. Bloomfield, O'Hara, and Saar (2005), in a laboratory experiment, find that informed traders submit more limit orders than market orders. They exploit their informational advantage early in the trading period to find mispriced limit orders moving the market towards the true price, thereby progressively reducing the value of their information. As the end of the trading period approaches, they switch increasingly to limit orders, as the value of their informational advantage falls away.

A number of papers have also attempted to test the informativeness of limit order book information empirically for different asset classes. We can classify these into two main groups: studies concerning the statistical significance of information contained in limit orders to explain future returns and studies that look at the economic value of exploiting this information. Within the first group, Cao, Hansch, and Wang (2009) considered the information content of a limit order book behind the best bid and offer using data from the Australian Stock Exchange. They found that the contribution of the order book to price discovery is approximately 22%, while the rest of the variation in future returns comes from the best bid, offer, and transaction prices. They demonstrate that order imbalances between the demand and supply schedules along the book are statistically significantly related to future short-term returns. Harris and Panchapagesan (2005) also found that information on limit order book depth forecasts short-term changes in prices. Hillman and Salmon (2007) using FX tick data explore the information content of the limit order book to explain returns using variogram techniques, which involves no specific parametric model. They show clear in-sample ability to explain very short run movements in the USD/DM rate using a range of measures of order book structure. Hellström and Simonsen (2009), using a count data time series approach, find that there is informational value in the first levels of the bid- and ask-side of the order book. They also show that both the change and the imbalance of the order book statistically significantly explain future price changes. Offered quantities at the best bid and ask prices on data from the Swedish Stock Exchange reveal more information about future short run returns than measures capturing the quantities at prices below and above. The impacts are most apparent at the one minute aggregation level, while results for higher aggregation levels generally show insignificant results. These results would suggest that the informational content of the order book is very short-term.

While the above-mentioned papers focus on in-sample statistical significance, there are also several papers demonstrating some ability of conditioning information to predict

future movements of returns out-of-sample.³ Huang and Stoll (1994) found that differences in quoted depth predict future returns at five-minute intervals out-of-sample. Evans and Lyons (2005, 2006) were among first to document the forecasting power of customer order flow to outperform a random walk benchmark. Froot and Ramadorai (2005) report that order flow contains some information for future exchange rate returns in low frequency data. Rime, Sarno, and Sojli, (2010) employ data for three major exchange rates from the Reuters electronic interdealer trading platform and confirm these findings. In contrast to the above studies, Danielsson, Luo, and Payne (2002) find limited and Sager and Taylor (2008) find no evidence of superior forecasting ability of order flow over random walk models at different forecast horizons.

There are not many papers falling into the second group that focus on whether or not the market information can be exploited economically by market participants. Chordia and Subrahmanyam (2004) find profitability of future stock returns using order imbalance; Della Corte, Sarno, and Tsiakas (2009) and Rime, Sarno, and Sojli (2010) find profitability of exchange rate returns using transaction order flow but only in the long run. However, neither of these papers consider limit order book information in their trading strategies. A notable exception is a paper by Latza and Payne (2010), who considered the forecasting power of market and limit order flows on stock returns and show that both can forecast returns. They show, via simulation, that dealers who time the execution of the trades on the limit order flow can reduce the cost of trading customer orders by up to 20%.

3. Data and methodology

We use interdealer tick-by-tick data for the U.S. dollar sterling exchange rate drawn from the Reuters D3000 trading system, which is the electronic broker trading platform where most sterling trades take place. In order to make sure that our results are not driven by the use of any particular sample period, we use five different data sets: weeks commencing on January 13, 2003, February 10, 2003, March 17, 2003, and two days on March 31 and April 1, 2008 (there is a much higher frequency of trades in 2008, as we discuss below). The data we analyze consists of continuously recorded limit and market orders and their volumes between 07:00 and 17:00 GMT which allows us to reconstruct the full limit order book on a tick-by-tick basis. For each entry, the data set contains a unique order identifier, quoted price, order quantity, quantity traded, order type, transaction identifier of order entered or removed, status of market order, entry type of orders, removal reason, and date and time of orders entered and removed. The data time stamp's precision is 1/100th of a second and the minimum trade size in Reuters electronic trading system is 1 million pounds sterling.

3.1. Summary statistics

Table 1 reports the summary statistics for transaction and limit order data for the different sample periods. Average inter-quote durations (speed of limit order arrival or

³This distinction is important both theoretically and empirically. It is widely recognized that in-sample fit does not necessarily translate to out-of-sample predictability. There is a range of reasons why this may be the case; in-sample overfitting to restricted sample information, model misspecification and structural changes are the main explanations provided in the literature. The wider issue concerns the fact that all inference is conditional in effect on the sample information, which may poorly represent the range of behavior in the full population (see Hansen, 2010 for a full discussion).

Table 1

Descriptive statistics on market liquidity.

The table presents summary statistics on the liquidity of the market for the five subsamples. It reports the mean, standard deviation, median, minimum, maximum, and first and the third quartiles of best quantities, slopes of bid and ask sides of limit order book, the depth, inter-quote duration, and bid–ask spread. Subsamples are: January 13–17, 2003, February 10–14, 2003, March 17–21, 2003, March 31, 2008, and April 1, 2008. Best quantities and depth are measured in millions of pounds sterling, slopes are basis point per 100 million of currency trade, duration is in seconds and bid–ask spread is in basis points.

Variable	January 13–17, 2003					February 10–14, 2003				
	Mean	Std. Dev.	Q1	Median	Q3	Mean	Std. Dev.	Q1	Median	Q3
Best bid quantity	2.80	2.53	1	2	3	3.07	5.64	1	2	3
Best ask quantity	2.90	7.45	1	2	3	3.03	5.94	1	2	3
Slope of bid side	58.69	77.87	25.92	40.00	66.67	55.53	52.43	25.00	40.00	66.67
Slope of ask side	65.35	66.08	28.57	50.00	75.00	66.71	77.11	27.77	44.44	75.00
Depth of bid side	44.43	18.45	29	44	58	48.81	34.46	26	39	56
Depth of ask side	39.63	25.86	23	32	46	34.21	14.30	25	33	42
Inter-quote dur.	2.29	16.75	0.30	1.00	2.40	1.82	3.26	0.30	0.95	2.11
Bid–ask spread	2.28	2.06	1	2	3	2.31	2.01	1	2	3
Nr. of orders			75,135					98,785		

Variable	March 17–21, 2003				
	Mean	Std. Dev.	Q1	Median	Q3
Best bid quantity	2.76	2.80	1	2	3
Best ask quantity	2.84	3.19	1	2	3
Slope of bid side	75.54	101.9	30.76	50.00	83.33
Slope of ask side	85.03	98.19	33.33	57.14	100.0
Depth of bid side	40.17	15.54	28	40	51
Depth of ask side	38.22	18.58	25	33	48
Inter-quote dur.	1.84	3.33	0.31	0.92	2.22
Bid–ask spread	2.73	3.16	1	2	3
Nr. of orders			97,559		

Variable	March 31, 2008					April 1, 2008				
	Mean	Std. Dev.	Q1	Median	Q3	Mean	Std. Dev.	Q1	Median	Q3
Best bid quantity	3.93	4.06	2	3	5	4.00	4.31	2	3	5
Best ask quantity	3.80	4.63	2	3	5	4.01	6.23	2	3	5
Slope of bid side	20.89	14.54	12.50	17.64	25.00	19.52	12.05	11.76	16.66	25.00
Slope of ask side	21.52	14.31	13.33	18.18	25.00	20.26	13.50	12.50	16.66	25.00
Depth of bid side	79.87	27.69	61	78	94	89.56	32.82	70	85	100
Depth of ask side	72.62	33.12	48	67	89	80.64	32.60	55	79	98
Inter-quote dur.	0.10	0.32	0.01	0.03	0.09	0.09	0.28	0.01	0.03	0.08
Bid–ask spread	2.53	1.21	2	2	3	2.49	1.22	2	2	3
Nr. of orders			594,519					388,259		

removal) are 2.29, 1.82, and 1.84 seconds for first three samples and 0.1 and 0.09 seconds for the two 2008 samples. This demonstrates that the electronic market is very active and critically its activity has grown tremendously from 2003 to 2008. There are 75,135, 98,785, and 97,559 orders for the three weeks in 2003 and 594,519 and 388,259 on March 31 and April 1, 2008, respectively. The average values of the bid–ask spread from our samples are 2.28, 2.21, and 2.73 basis points in 2003 and 2.53 and 2.59 basis points in 2008, indicating that D3000 is a very tight market.

Although there is no major difference in the bid–ask spreads between the 2003 and 2008 samples, there is a huge jump in market liquidity as measured by the slope of the limit order book and its depth. The average slopes of the limit order book in the 2003 subsamples are 58.69, 55.53, and 75.54 basis points per billion of currency trade for the bid side and 65.35, 66.71, and 85.03 for the ask side.⁴ In the 2008 samples, the slope values were 20.89 and 19.52 for bid side and 21.52 and 20.26 for ask side, indicating that the limit order book became about three times flatter in 2008 than it was in 2003. Also, the depth of the market almost doubled. These summary statistics indicate that the currency pair we are studying is traded in a highly liquid market.

3.2. Hypotheses

We are interested in two main hypotheses. The first is whether the exchange rate is predictable in terms of statistically significant economic value (economic predictability). Thus,

Hypothesis 1. The exchange rate returns is not economically predictable at high-frequency.

The second question is whether limit order book information adds economic value over that provided by the basic price information available.

Hypothesis 2. Limit order book information does not add significant economic value to the predictability of the exchange rate at high-frequency.

As mentioned above, we build trading strategies that are designed to exploit any profitable pattern in the exchange rate and test the profits obtained for significance. By varying the information set used in these trading rules, we are able to differentiate the predictive power of limit order book information from that contained in past prices and volumes.

As the majority of existing research has focused on linear predictive models, we also employ a *linear model* to forecast future exchange rate movements as a benchmark. Apart from this linear model we also use a *genetic algorithm* as a general non-parametric device to construct trading rules. This approach has the advantage that it is *model free* and designed to exploit both linear and any *non-linear* dependency between future returns and predictors. The specification of the genetic algorithm trading rules evolve according to their “fitness”, which is determined by an economic profit-based criterion. This approach

⁴We construct the slope of the demand and supply curves in our limit order book using the two best bid and ask quotes and the associated depth at these quotes. See also Section 3.3 for detailed description of variables.

is not therefore susceptible to the criticism that any result we find would have been due to the assumption of a specific trading rule we had selected *ex ante*.

Since we want to keep our trading strategies implementable in “real time”, we need to ensure that they are based exclusively on historical data available at the time of trade. We use an in-sample period to construct the rules and then check their performance out-of-sample. We describe the implementation of our approach next.

3.3. The information sets

We define four different conditioning information sets to test the added value of various limit order book variables. The values of these variables serve as inputs to a function (either the linear rule or the genetic algorithm) generating the trading signal.

1. *Screen information* (denoted by *Screen* hereafter) contains best limit order prices (both bid and ask) and their quantities as time series, the bid–ask spread, the level of mid-quotes, and the inter-quote duration. This information is considered as the basic set, which is normally available to all traders. It is also contained as a subset in the three information sets defined below.
2. *Limit order book information* (denoted by *Book* hereafter); in addition to the variables mentioned above, includes total depth, the number of layers in the limit order book, the difference between the best and the second best price (both, bid and ask), the slopes of the bid and ask curves of the limit order book, time series of the levels of quantity weighted quotes, and the quantity weighted mid-quote, the quantity weighted bid–ask spread, and the difference between the mid-quote and the quantity weighted mid-quote. By depth we mean the total quantity available at the moment in the limit order book on the particular side of the book (demand or supply). The slope of the bid side of the limit order book is defined as

$$\text{slope}^{bid} = (p_1^{bid} - p_2^{bid}) / q_1^{bid},$$

where p_1^{bid} and p_2^{bid} are the best and the second best prices on the bid side respectively and q_1^{bid} is the quantity available at the best bid price. The slope of the ask schedule is defined analogously. The quantity weighted bid price is defined as

$$wp^{bid} = \left(\sum_i (p_i^{bid} \times q_i^{bid}) \right) / \sum_i q_i^{bid},$$

where the index i runs through all available levels of bid quotes. The quantity weighted mid-quote is

$$wmid = (wp^{bid} + wp^{ask}) / 2$$

and the quantity weighted bid–ask spread is

$$wspread = wp^{ask} - wp^{bid}.$$

3. *Order flow information* (denoted by *Order* hereafter) contains the screen information set plus order flow information. Following [Latza and Payne \(2010\)](#), we use two different types of order flow: limit order flow and transaction order flow. Limit order flow is further decomposed into order flow on the best prices (the inside order flow) and order flow outside the best prices (the outside order flow). We construct 1 and 20 minute as well as 1 tick order flow variables for each side of the limit order book. By 1 tick order flow we mean the volume

of the most recent order of the corresponding type (either market order, limit order at the best price or the limit order outside the best price) strictly preceding the time of decision making.

4. *Full information* (denoted by *Full* hereafter) combines all three types of conditioning information mentioned above.

The Reuters D3000 trading platform did not allow traders to see limit orders and their quantities outside the best bid and ask prices in 2003 and 2008. Hence it would be impossible for traders to build strategies that explicitly use this information. Thus, if the three extended information sets are found to convey substantial information that can be profitably exploited, we will detect it in our experiments. If we cannot find added value, then it would imply that limit orders placed outside the best prices do not carry any significant information about future returns.

3.4. The trading mechanism and fitness function

We measure the fitness of trading rules by means of the cumulative returns from the following simple trading strategy. The trader buys or sells short 1 million pounds sterling according to the signal provided by the selected trading rule. This allows us to control for the potential price impact of trade since we can ensure that the liquidity necessary to complete a transaction with the minimum trade size is present in the market.⁵ As new information arrives from the market, the trader re-evaluates the trading signal and updates his position accordingly. This means that as soon as the trader observes any change in the limit order book, he can change or keep the same position depending on the outcome of the signal.

Under such a trading scheme, the trader is potentially able to trade at every single instant. In order to control for trading frequency, we add a trading threshold to the strategy. According to this, the trader is allowed to trade only if the exchange rate exceeds a band of $\pm k$, relative to his last transaction price. More formally, let z_t denote the state of the investor's position at time t . That is, $z_t = 1$ corresponds to a long position in sterling and $z_t = -1$ corresponds to a short position. The trader will re-evaluate his position only if $|p_t - p_{t_1}| \geq k$, where p_t is price at time t and t_1 denotes time of the trader's last transaction.

The coefficient k serves as an inertia parameter to filter out weak trading signals. The idea of such "filter rules" goes back to Alexander (1961) and Fama and Blume (1966). The parameter k determines an "inertia band" that prompts one to trade only once a realization of the exchange rate exceeds the value of a certain characteristic (past realized values of the exchange rate in our case) by a value of k . A larger inertia band (larger k) filters out more trades, thus reducing trading frequency. The use of an inertia parameter also has a behavioural interpretation based on the notion of ambiguity aversion. For instance, Easley and O'Hara (2010) show that in the face of Knightian uncertainty incomplete preferences may lead to an absence of trading. Traders will revise their position only if the trading signal is confirmed by other criteria that they have at their disposal, which is very often provided by simple technical tools.

Table 2 presents how many times the trader re-evaluates the position during the out-of-sample period for different values of k . A zero value of k means that trades can take place every time the mid-quote of the exchange rate changes and hence exhibits the largest

⁵We assume that traders can execute transactions at the current price immediately and are not affected by either latency problems nor execution risk.

Table 2
Number of indicative trades.

This table provides the number of times the mid-quote of the exchange rate goes outside the k -band for different values of k . When the exchange rate crosses the k -band, trading rules re-evaluate positions taken in the exchange rate. Trade occurs when the trading rule requires a change in the direction of the position. k is measured in basis points.

k	January 13–17, 2003	February 10–14, 2003	March 17–21, 2003	March 31, 2008	April 1, 2008
0	16,889	22,468	19,654	19,263	12,005
2	1,745	2,775	1,366	1,255	744
4	497	877	443	421	221
6	252	425	215	218	108
8	144	251	140	128	71
10	97	166	99	83	39
12	69	117	64	58	
14	59	83	45	46	23
16	41	65	34	37	17
18	31	48	24	29	12
20	23	35	22	20	11
22	19	31	22	18	13
24	15	25	21	18	7
26	15	23	15	12	6
28	11	15	14	12	4
30	10	15	10	10	4

number of transactions. As k increases, the trading frequency drops. For $k=30$, only up to 10 transactions per day can be made.

We use simple cumulative returns as a performance measure to evaluate the profitability of trading strategies:

$$R_c = \prod_t (1 + z_t r_t) - 1,$$

where $r_t = (p_t - p_{t-1}) / p_{t-1}$ is the one-period return of the exchange rate. Here p_t denotes the corresponding best bid p_t^{bid} or best ask p_t^{ask} price.

3.5. The linear trading rule

The linear model's predictions are generated using a linear regression. We use an in-sample period to estimate the regression model where the dependent variable is the one step ahead mid-quote exchange rate return r_{t+1} and the regressors are time t dated values of all the variables contained in the relevant information set. Out-of-sample forecasts of future exchange rate returns serve as signals for a simple binary trading rule, i.e., positive (negative) predicted values of future returns are associated with a “buy” (sell) signal. Based on these signals, we construct a trading strategy as described above and evaluate its out-of-sample performance for different values of the inertia parameter k .

3.6. Genetic algorithm trading rule

The genetic algorithm provides an effective method for searching over space of potential trading rules, both linear and non-linear. This method allows us to evaluate predictability as generally as possible and not impose any effective restriction on the form of the model, predictor or trading rule. The genetic algorithm is a computer-based optimization procedure that uses the evolutionary principle – the survival of the fittest – to find an optimum. It provides a systematic search process directed by performance rather than gradient.⁶

Starting from an initial set of rules, the genetic algorithm evaluates the fitness of various candidate solutions (trading rules) using the given objective function. It provides as an output, solutions that have higher in-sample cumulative returns on average.

We build a trading rule as a binary logical tree, which produces true or false signals given the set of input variables. If the value of the rule is “true”, it gives the signal to “buy” an asset. If the rule is “false”,—the trader “sells” the asset short. The rules are represented in the form of randomly created binary trees with terminals and operations in their nodes. We employ the following choices of operations and terminals.

Operations: The function set used to define the technical rules consists of the binary algebraic operations $\{+, -, *, /, \max, \min\}$, binary order relations $\{<, >, \leq, \geq, =\}$, logical operations $\{\text{and}, \text{or}, \text{and}\}$ and unary functions $\{\text{abs}, -\}$ of absolute value and change of sign.

Terminals: The terminal set contains the variables, which take their values from data and are updated every time new information arrives in the market. Thus, it allows the conditioning information sets to update the trading rule as time passes. The genetic algorithm also explicitly computes lag values of the conditioning variables, their moving average values, and maxima and minima over different periods. The terminal set also includes real numbers as terminal constants.

An example of a tree and the corresponding trading strategy is given in Fig. 1. It presents a trading rule that generates a signal “buy” if the current quantity weighted spread is less than the last 10 trades average returns times the bid–ask spread and the absolute value of the difference between the quantity weighted bid and best bid price is less than 0.001. Otherwise, the signal is “sell”.

Two operations of crossover and mutation are applied to create a new generation of decision rules based on the genetic information of the fittest candidate solutions.

Crossover: For the crossover operation, one randomly selects two parents from the population based on their fitness. A node within each parent is then taken as a crossover point selected randomly and the subtrees at the selected nodes are exchanged to generate two children. One of the offspring then replaces the less fit parent in the population. In our implementation, we use a crossover rate of 0.4 for all individuals in the population. This operation combines the features of two parent chromosomes to form two similar offspring by swapping corresponding segments of the parents. In our case, these segments are

⁶Nix and Vose (1992) and Vose (1993) use a Markov Chain framework to show that asymptotically in population size – populations that have suboptimal average fitness have probabilities approaching zero in the stationary distribution of the Markov Chain, whereas the probability for the population that has optimal average fitness approaches one. The genetic algorithm’s success as an optimizer depends on having a sufficiently large population of individual rules. We have taken considerable effort to ensure that this is the case in our experiments with the results repeated below representing only a fraction of a very large computational exercise.

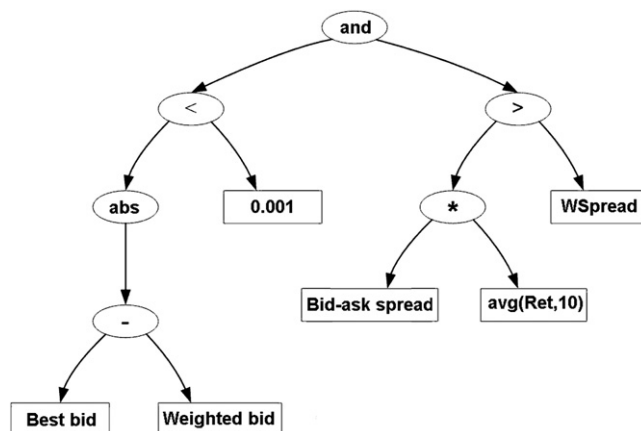


Fig. 1. Example of genetic algorithm trading rule.

The figure presents an example of a genetic algorithm-based trading rule in a form of a binary logical tree. It generates a “buy” signal if the current quantity weighted spread is less than the past ten periods average returns times the bid–ask spread and the absolute value of difference between the quantity weighted bid and the best bid price is less than 0.001. Otherwise, the signal is “sell.”

represented by sub-nodes of a binary tree. The intuition behind the crossover operator is information exchange between different potential solutions.

Mutation: In order to mutate a rule, one of its subtrees is selected at random and replaced with the new randomly generated tree. This operation guarantees the refreshment of the genetic code within the population. The best 25% of the rules are not mutated at all and the remaining are mutated with probability 0.1. The intuition behind the mutation operator is the introduction of some extra variability into the population of trading rules.

The evolutionary algorithm can be summarized as follows:

1. Create randomly the initial population $P(0)$ of trading rules given the information set and initialize the number of iterations $i=0$;
2. Set $i := i + 1$;
3. Evaluate in-sample fitness of each tree in the population using the fitness function;
4. Generate a new population of trees (i.e., the set of new trading rules) using the genetic operations (crossover and mutation) and replace the old population with the new one;
5. Repeat 2–5 while $i < N$.

After each such iteration, rules that have poor performance according to the fitness function are removed from the population and only the more profitable candidates survive and carry their structure onwards to create new trading rules. Ultimately, the algorithm converges to the trading rule achieving the best in-sample performance given the conditioning information.

In the program we have experimented and use a population size of 200 individual trading rules and perform 1,000 iterations of the algorithm (that is, $N=1,000$).

The complexity of trading rules is controlled in a probabilistic manner. In fact, the probability for a binary node to appear in the tree is smaller than the probability of a unary one, which prevents the tree from becoming very large.

3.7. Testing procedures

We employ a series of statistical tests to test for profitability of each of the trading strategies. We split each trading period into two equal parts that serve as in-sample and out-of-sample periods respectively. We use the in-sample period as the estimation sample for the linear regression model. We test the economic value of a strategy using Anatolyev-Gerko statistic (Anatolyev and Gerko, 2005). This test compares the profitability of a trading strategy relative to the random walk model. The relative performance of the trading strategies are based on different conditioning information sets and then tested using the Giacomini-White test for conditional predictive ability (Giacomini and White, 2006).

Similarly, with the genetic algorithm-based strategy, we choose the trading rule that produces the best in-sample performance and test its profitability out-of-sample. In order to generate an empirical distribution of the out-of-sample cumulative returns, we run this procedure independently 100 times. This provides us with potentially (due to the stochastic nature of the genetic algorithm search) 100 different trading rules and their out-of-sample performance. Using this sample of independent cumulative out-of-sample returns, we can use a t -statistic to test if the mean of the returns is significantly different from zero. The relative performance of the different information sets is tested using a paired t -test. Specifically we test if the difference in the unconditional mean of returns for two strategies is based on different information sets that are significantly different from zero.

In addition, we combine signals from the 100 best individual trading rules and create an aggregated genetic algorithm signal, which we call the “majority” rule. This combined rule is an alternative strategy to the single best in-sample genetic algorithm rule. It produces a “buy” (sell) signal if the majority of the 99 independent best in-sample rules produce the “buy” (sell) signal. This rule probably reflects the way in which technical analysis is used by practitioners. Traders often do not follow a single rule but form an impression as to where the market is moving on the basis of a number of technical indicators, dropping those that appear not to have worked well in the past. The economic value of this rule is then tested using the Anatolyev-Gerko test and the relative performance is tested by the Giacomini-White test.

We carry out our exercises by allowing the trader to trade using bid and ask prices (taking into account transaction costs explicitly).⁷ The trader always buys at the best ask price and sells at the best bid price, so the current bid–ask spread reflects the real transaction costs a trader would face in the market.

4. Results

We start by examining Hypothesis 1 as to whether there is evidence for the predictability and profitability of exchange rates at high frequency and then we consider the relative performance of the different information sets.

4.1. Hypothesis 1: predictability and profitability

Coefficients estimates for the different information sets for the linear model are given in Table 3. The results show that for each of the three extended information sets beyond

⁷The corresponding results in the case of no transaction costs are qualitatively similar and are available upon request.

Table 3

Linear model parameters estimates.

This table presents coefficient estimates from the linear regression of one-step ahead returns on different conditioning variables. The estimation period is the first half of the week commencing on January 13, 2003. The “Screen” column contains estimates of the variables from the *Screen* information set, “Book” corresponds to the limit order book information set, “Order” denotes order flow information and “Full” provides results for the combined information set. The table does not include variables that are not significant in any of the information sets. The *F*-test row presents *F*-statistic values for testing the joint significance of all variables added in addition to the *Screen* information set variables (in the “Screen” column the usual *F*-statistic for significance of the regression is given). * and ** indicate significance at the 5% and 10% levels.

Variable	Screen	Book	Order	Full
Intercept	$-4.43 \times 10^{-6**}$	$-4.66 \times 10^{-6*}$	$-5.50 \times 10^{-6**}$	$-6.93 \times 10^{-6**}$
Best ask quantity	$-5.30 \times 10^{-7**}$	$-5.28 \times 10^{-7**}$	$-5.39 \times 10^{-7**}$	$-5.63 \times 10^{-7**}$
Best bid quantity	$1.12 \times 10^{-6**}$	$9.33 \times 10^{-7**}$	$1.01 \times 10^{-6**}$	$8.66 \times 10^{-7**}$
Ret_{t-1}	$-0.296**$	$-0.295**$	$-0.312**$	$-0.311**$
Inter-quote duration	5.02×10^{-8}	7.42×10^{-8}	9.62×10^{-8}	1.08×10^{-7}
Bid-ask spread	$0.0116**$	$0.0129**$	$0.0117**$	$0.0129**$
Slope of ask side		$1.08 \times 10^{-4**}$		$1.17 \times 10^{-4**}$
Slope of bid side		$-2.20 \times 10^{-4**}$		$-2.30 \times 10^{-4**}$
Depth of ask side		$5.69 \times 10^{-8*}$		5.16×10^{-8}
Mid-quote difference		$0.00121**$		$0.00132**$
Quantity weighted spread		2.88×10^{-4}		2.77×10^{-4}
Best bid order flow(1 tick)			$1.48 \times 10^{-6**}$	$1.50 \times 10^{-6**}$
Bid cancel order flow(1 tick)			$-8.81 \times 10^{-7**}$	$-8.83 \times 10^{-7**}$
Bid cancel order flow(1 min)			$8.64 \times 10^{-8*}$	$1.08 \times 10^{-7**}$
Best ask order flow(1 tick)			$-1.55 \times 10^{-6**}$	$-1.51 \times 10^{-6**}$
Best ask order flow(20 min)			2.66×10^{-8}	$5.78 \times 10^{-8**}$
Ask order flow(20 min)			-1.87×10^{-8}	$-4.65 \times 10^{-8**}$
Ask cancel order flow(1 tick)			$5.87 \times 10^{-7*}$	$6.00 \times 10^{-7**}$
Ask cancel order flow(1 min)			-1.09×10^{-7}	$-1.31 \times 10^{-7**}$
Transaction order flow(1 tick)			$1.60 \times 10^{-6**}$	$1.43 \times 10^{-6**}$
<i>F</i> -test	1176.13	10.96	5.46	7.13

Screen information, the majority of the explanatory variables are statistically significant in-sample. Moreover, we can easily reject the joint hypothesis that all coefficients in each of the extended information sets beyond the *Screen* information set are insignificant from zero on the basis of an *F*-test. This confirms the results in the existing literature that statistically, limit order book information does contribute to the in-sample explanation of the exchange rate. In order to check whether this apparent predictability can be translated into out-of-sample profitability, we implement the trading strategy described above using the in-sample coefficients estimates.

Table 4 reports the average out-of-sample returns of the genetic algorithm trading rules for different values of *k* under no transaction costs for the five sample periods.⁸ Using the empirical distribution of generated from the out-of-sample performance of the best 100 genetic algorithm rules, we use a *t*-statistic to test if the average return is statistically different from zero. The trading strategies produce high positive average daily returns for

⁸All returns presented in this paper are adjusted on a daily basis.

Table 4

Out-of-sample performance of the genetic algorithm without transaction costs.

This table presents the out-of-sample cumulative returns (in percent) generated by the genetic algorithm for the five different sample periods. The in-sample period (the first half of the sample) is then used to select the best performing trading rule and this rule is used to trade out-of-sample (the second half of the sample). The exercise is repeated 100 times to generate the empirical distribution of cumulative returns. Returns are adjusted to a daily basis and expressed in percent. k is the threshold value for the trading band—a trader re-evaluates his position only if the current mid-quote is outside the k -band around the last traded price. k is measured in basis points. There are no transaction costs involved and trading is based solely on mid-quotes returns. Columns contain the results for the different information sets. “Screen” denotes the *Screen* information, “Book” denotes the limit order book information set, “Order” denotes the order flow set and “Full” corresponds to the combined information set. * indicates 5% significance level according to t -test statistics.

k	January 13–17, 2003				February 10–14, 2003				March 17–21, 2003			
	Screen	Book	Order	Full	Screen	Book	Order	Full	Screen	Book	Order	Full
0	11.01*	10.67*	10.61*	10.13*	14.88*	14.53*	14.43*	13.72*	17.19*	15.68*	15.74*	15.07*
2	3.69*	3.59*	3.26*	3.22*	5.94*	5.26*	5.25*	5.03*	7.05*	6.53*	6.68*	6.50*
4	1.12*	1.22*	1.07*	1.09*	1.38*	1.59*	1.33*	1.53*	0.20	0.67*	0.59*	0.75*
6	0.43*	0.44*	0.01	0.29*	0.43*	0.48*	0.09*	0.34*	0.21*	0.68*	0.11*	0.59*
8	0.07	0.16*	0.14*	0.07	0.66*	0.63*	0.39*	0.55*	0.07	0.32*	−0.20*	0.00
10	0.15*	0.16*	0.18*	0.21*	0.58*	0.70*	0.36*	0.52*	−0.12*	−0.31*	−0.20*	−0.20*
12	−0.04*	0.08*	0.02	0.07	0.45*	0.72*	0.17*	0.51*	−0.11*	−0.08*	0.06*	0.01
14	−0.12*	0.19*	−0.07	0.15*	0.80*	0.84*	0.62*	0.77*	0.12*	0.08*	0.02	0.15*
16	−0.23*	0.28*	−0.17*	0.14*	−0.09	0.32*	−0.28*	0.08	0.12*	0.01	0.05	−0.05
18	−0.04	0.03	−0.14*	0.04	0.20*	0.18*	−0.10*	0.08*	0.15*	0.00	0.13*	0.07
20	−0.30*	−0.24*	0.02	−0.10*	−0.07*	0.22*	−0.10*	0.06	0.01	−0.06*	0.02	0.04
22	−0.15*	−0.16*	−0.06*	−0.09*	0.00	−0.05	0.01	0.05*	0.16*	0.08*	0.09*	0.16*
24	−0.11*	−0.12*	−0.01	−0.05	−0.08*	−0.33*	−0.03	−0.25*	0.32*	0.08*	0.06	0.06
26	−0.18*	−0.17*	−0.04	−0.04	0.03	−0.02	−0.19*	−0.13*	−0.17*	−0.11*	−0.07*	−0.13*
28	−0.01	−0.01	0.06	0.01	−0.04	0.07*	0.04	0.06*	−0.14*	−0.02	−0.01	−0.11*
30	−0.09*	−0.10*	−0.06*	−0.08*	0.03	0.02	0.06*	0.06*	−0.16*	−0.18*	−0.12*	−0.18*

k	March 31, 2008				April 1, 2008			
	Screen	Book	Order	Full	Screen	Book	Order	Full
0	29.01*	26.74*	24.59*	26.98*	19.13*	17.05*	16.56*	17.47*
2	2.09*	2.52*	2.68*	1.65*	1.41*	1.10*	−0.01	−0.36*
4	0.07*	0.13	0.51*	0.11	−0.84*	−0.62*	−0.57*	−0.76*
6	0.03	−0.03	0.16*	0.05	−0.25*	−0.70*	−0.69*	−0.54*
8	1.46	1.41*	1.06*	0.94	0.08	0.08	0.11	0.05
10	0.47*	0.26*	0.30*	0.49*	0.33*	0.40*	0.17	0.17
12	0.28*	0.13	0.22*	0.22*	−0.75*	−0.23*	−0.50*	−0.52*
14	0.38*	0.71*	0.03	−0.32*	0.31*	−0.20*	0.10	0.13
16	1.18*	1.34*	0.71*	0.62*	−0.59*	−0.41*	−0.47*	−0.39*
18	0.44*	0.62*	0.28*	0.48*	−0.47*	−0.36*	0.43*	0.10
20	−0.99*	−0.70*	−0.68	−0.71*	−0.43*	−0.52*	0.02	0.13
22	0.30*	−0.09	−0.23*	−0.43*	0.00	−0.44*	−0.60*	−0.46*
24	−0.99*	−0.14	−0.85	−1.03*	−0.45*	−0.44*	−0.29*	−0.39*
26	−0.26*	−0.17	−0.23	−0.13	−0.57*	−0.50*	−0.57*	−0.57*
28	−0.28*	−0.32*	−0.31	−0.35*	−0.78*	−0.41*	−0.51*	−0.71*
30	−0.06	0.42*	0.08	0.20*	0.04	−0.39*	−0.59*	−0.79*

small values of the inertia parameter k (up to 6 basis points for the 2003 samples, up to 4 basis points for the March 31, 2008 sample and only for $k=0$ for the April 1, 2008 sample). The t -test indicates that the average returns are significantly different from 0 in these cases. Profitability disappears as the frequency of trading decreases. In fact, after $k=12$, the trading rules start generating negative out-of-sample returns during January 13–17, 2003, for $k=16$ during February 10–14, 2003 and for $k=8$ during March 17–21, 2003. For the 2008 samples negative out-of-sample returns start to appear at $k=6$ for the March 31 sample and at $k=2$ for the April 1 sample.

A more interesting question is whether the profitability of the trading strategies remains when transaction costs are incorporated.⁹ Table 5 reports the out-of-sample performance of the linear model when trading on the best bid and ask prices and tests for superior predictability relative to a random walk using the Anatolyev-Gerko test.

Table 5 shows that returns drop substantially when trading at a high frequency. Small values of the inertia parameter reflect a higher number of transactions, which implies a large cumulative transaction cost that exceeds the profits from trading. As the number of transactions drops, transaction expenses decrease and trading rules become profitable again.¹⁰ Out-of-sample returns are positive at k ranges from 6 to 16 basis points during January 13–17, 2003, at k from 8 to 28 basis points during February 10–14, 2003 and k from 10 to 22 basis points during March 17–21, 2003. There are less pronounced patterns of positive returns for the 2008 sample. On March 31, only k values of 10, 20, and 22 basis points generate positive returns; there are positive returns on April 1 for k ranging from 8 to 12 and from 16 to 20 basis points. However, it is important to note that for all samples at $k=0$ and 2, the linear trading rule is superior to the random walk across all information sets. Although it generates negative returns, the results indicate that the random walk model loses even more based on trading with transaction costs. For each sample there are inertia values within the 10–14 basis points range that generate positive out-of-sample returns superior to the random walk model (with the exception for the March 31, 2008 sample for the *Full* information set).

The highest daily returns are 0.72% during January 13–17, 2003 achieved at $k=12$, 1.79% during February 10–14, 2003 at $k=14$, 0.89% during March 17–21, 2003 at $k=20$, 1.93% during March 31, 2008 and 2.46% on April 1, at $k=12$. Although these returns may look very high, there is no obvious pattern as to how to exploit them systematically due to the changing nature of the optimal inertia parameter. We address this question at the end of the section. Also, note that traders cannot invest any desired amount of capital into the trading strategy due to our restriction on the trade size.

Importantly, Table 6 shows using the t -test for zero mean that the genetic algorithm can handle transaction costs surprisingly well. While the inertia parameter is the only way to mediate the trading frequency for the linear trading rule, the genetic algorithm can adjust it endogenously.¹¹ The trading strategy based on the genetic algorithm shows positive and significant positive returns under transaction costs even for low k bands for the 2003 data. Profitability vanishes though for values of k higher than 12 basis points. However, there is

⁹All further results presented in the paper take transaction costs into account.

¹⁰This result is in line with the findings of Knez and Ready (1996), Cooper (1999), and Balvers and Wu (2010), who found that the after-transaction-cost returns with “filter rules” improve compared to trading strategies with a zero filter.

¹¹Note that genetic algorithm produces different trading rules for trading with no transaction costs and with transaction costs.

Table 5

Out-of-sample performance of the linear model including transaction costs.

This table presents the out-of-sample cumulative return generated by the linear model for the five different sample periods. Linear regression coefficients are estimated using in-sample data (first half of the sample) and returns predictions are formed for the out-of-sample period (the second half of the sample). A simple binary trading rule is implemented based on the return predictions. Returns are adjusted to a daily basis and expressed in percent. k is the threshold value for the trading band—a trader re-evaluates his position only if the current mid-quote is outside the k -band around the last traded price. k is measured in basis points. Transaction costs are reflected in bid–ask spread as trading is based on best bid and ask limit orders. Columns contain results for the different information sets. “Screen” denotes the *Screen* information, “Book” denotes the limit order book information set, “Order” denotes the order flow set and “Full” corresponds to the combined information set. * indicates significance at the 5% level according to Anatolyev-Gerko test statistics.

k	13–17 January 2003				10–14 February 2003				17–21 March 2003			
	Screen	Book	Order	Full	Screen	Book	Order	Full	Screen	Book	Order	Full
0	–37.1*	–36.9*	–36.9*	–36.7*	–47.4*	–46.7*	–47.5*	–46.3*	–53.7*	–53.5*	–53.8*	–53.6*
2	–4.34*	–4.18*	–4.22*	–4.17*	–5.69*	–5.68*	–5.75*	–5.52*	–10.9*	–10.9*	–10.9*	–11.0*
4	–0.10	–0.17	–0.13	–0.18	–0.78	–0.87*	–0.73	–0.74	–2.37*	–2.39*	–2.31*	–2.36*
6	0.52	0.56	0.55	0.58	–0.28	–0.36	–0.29	–0.38	–0.80*	–0.73*	–0.66	–0.63
8	0.31	0.25	0.19	0.22	0.41	0.52	0.23	0.46	–0.35	–0.28	–0.40	–0.27
10	0.55*	0.55*	0.55*	0.60*	0.90*	1.00*	0.85*	0.93*	–0.35	–0.31	–0.38	–0.38
12	0.72*	0.72*	0.62*	0.61*	1.20*	1.09*	1.20*	1.09*	0.33	0.31	0.33	0.36
14	0.71*	0.71*	0.71*	0.71*	1.79*	1.71*	1.79*	1.71*	0.72*	0.73*	0.63*	0.64*
16	0.19	0.19	0.19	0.19	1.48*	1.56*	1.48*	1.56*	0.27	0.27	0.34	0.32
18	–0.02	–0.02	0.07	0.07	1.12*	1.12*	1.04*	1.04*	0.51	0.51	0.42	0.42
20	–0.39	–0.39	–0.29	–0.29	1.43*	1.54*	1.43*	1.65*	0.89*	1.00*	0.91*	1.11*
22	–0.04	–0.04	–0.04	–0.04	0.74*	0.74*	0.74*	0.74*	0.88*	1.01*	0.88*	1.01*
24	–0.26	–0.26	–0.26	–0.26	0.79*	0.79*	0.66*	0.54*	0.64*	0.64*	0.64*	0.54
26	–0.11	–0.11	–0.11	–0.11	0.54*	0.43*	0.54*	0.29	–0.10	–0.10	0.03	0.03
28	–0.25	–0.25	–0.25	–0.25	0.09	0.09	0.09	0.09	0.19	0.19	0.19	0.19
30	–0.20	–0.20	–0.20	–0.20	–0.01	0.16	0.01	0.04	–0.13	0.23	–0.13	0.23

k	March 31, 2008				April 1, 2008			
	Screen	Book	Order	Full	Screen	Book	Order	Full
0	–92.4*	–92.4*	–83.1*	–88.4*	–80.1*	–80.2*	–79.7*	–78.4*
2	–13.8*	–14.0*	–8.92*	–11.8*	–6.48*	–6.50*	–6.09*	–5.97*
4	–4.72	–4.74	–2.03	–5.72	–1.52*	–1.58*	–1.22	–0.88*
6	–0.55	–0.72	–1.26	–1.51	–0.02	–0.02	–0.20*	–1.32
8	0.60	0.72	0.73	1.60	1.49*	1.49*	1.29	0.63
10	0.45*	0.45*	–0.04*	–0.36*	1.84*	1.84*	1.97*	1.78*
12	0.46*	0.46*	1.61*	–0.51*	2.46*	2.46*	1.23	0.27*
14	0.32*	0.32*	0.84*	–0.53*	1.16*	1.16*	–0.03	–0.03
16	0.32	0.60	–1.12	–0.17	0.90	0.90	0.51	0.51
18	1.56	1.56	1.76	1.93	0.21*	0.59*	0.53	0.53
20	0.06	0.06	0.41	0.30	0.31*	0.31*	0.09	0.03
22	0.56	0.56	0.10	–0.89	1.25*	1.25*	–0.11	0.39
24	0.10	0.10	–0.29	0.54	0.17*	0.17*	–0.34	–0.34
26	–0.29	–0.29	–0.23	1.30	–0.46	–0.46	0.05	0.05
28	–0.16	–0.16	–1.27	–0.73	–0.44	–0.44	–0.44	0.09
30	–0.91	–0.91	–0.50	–1.47	–0.54	–0.54	–1.14	–1.14

Table 6

Out-of-sample performance of the genetic algorithm including transaction costs.

This table presents the out-of-sample cumulative returns (in percent) generated by the genetic algorithm for the five different sample periods. The in-sample period (the first half of the sample) is then used to select the best performing trading rule and this rule is used to trade out-of-sample (the second half of the sample). The exercise is repeated 100 times to generate the empirical distribution of cumulative returns. Returns are adjusted to a daily basis and expressed in percent. k is the threshold value for the trading band—a trader re-evaluates his position only if the current mid-quote is outside the k -band around the last traded price. k is measured in basis points. Transaction costs are reflected in bid-ask spread as trading is based on best bid and ask limit orders. Columns contain results for the different information sets. “Screen” denotes the Screen information, “Book” denotes the limit order book information set, “Order” denotes the order flow set and “Full” corresponds to the combined information set. * indicates significance at the 5% level according to t -test statistics.

k	January 13–17, 2003				February 10–14, 2003				March 17–21, 2003			
	Screen	Book	Order	Full	Screen	Book	Order	Full	Screen	Book	Order	Full
0	1.62*	1.60*	1.55*	1.51*	0.96*	0.97*	0.75*	0.78*	0.28*	0.06	−0.12*	−0.10*
2	1.50*	1.50*	1.41*	1.43*	0.97*	0.96*	0.94*	0.87*	0.39*	0.22*	0.08*	−0.07*
4	1.57*	1.39*	1.26*	1.30*	1.10*	1.12*	0.89*	0.84*	0.42*	0.31*	0.00	−0.01
6	0.87*	0.42*	0.20*	0.07	0.93*	0.72*	0.51*	0.19*	0.22*	0.09*	−0.11*	−0.18*
8	0.83*	0.61*	0.57*	0.45*	0.78*	0.61*	0.40*	0.38*	0.03	0.02	−0.36*	−0.34*
10	0.59*	0.41*	0.59*	0.41*	0.82*	0.76*	0.60*	0.77*	−0.17*	−0.14*	−0.28*	−0.31*
12	0.13*	0.34*	0.11	0.31*	1.03*	1.02*	0.58*	0.73*	0.11*	0.00	−0.01	−0.05
14	0.21*	0.48*	−0.01	0.25*	1.11*	1.13*	0.86*	1.08*	0.21*	0.15*	0.22*	0.19*
16	0.04	0.45*	−0.07	0.13*	0.51*	0.56*	−0.12*	0.31*	0.08*	0.05*	−0.04	−0.13*
18	−0.03	0.13*	−0.10*	−0.05	0.28*	0.23*	−0.21*	−0.05	0.19*	0.04	0.25*	0.17*
20	−0.10*	−0.09*	0.08*	0.12*	0.25*	0.39*	0.34*	0.33*	−0.08*	−0.10*	0.19*	0.10*
22	0.00	−0.09*	−0.03	−0.07*	−0.08*	0.10*	−0.02	0.02	0.19*	0.24*	0.18*	0.24*
24	−0.02	−0.06	0.06*	0.00	0.06	−0.35*	0.08*	−0.10*	0.39*	0.05	0.26*	0.16*
26	−0.14*	−0.10*	−0.04	−0.04	0.19*	0.06	−0.14*	−0.08*	−0.10*	−0.24*	−0.16*	−0.14*
28	0.01	−0.12*	0.03	−0.06	0.05	−0.04	0.07*	0.01	−0.03	−0.02	−0.04	−0.07*
30	−0.05	−0.03	−0.02	−0.03	−0.01	−0.04	0.00	−0.06*	−0.22*	−0.33*	−0.18*	−0.25*

k	March 31, 2008				April 1, 2008			
	Screen	Book	Order	Full	Screen	Book	Order	Full
0	0.88*	0.52*	0.45*	0.32*	0.46*	0.07	−0.45*	−0.21*
2	0.04	−0.12	−0.07	−0.05	−0.74*	−0.71*	−1.99*	−2.01*
4	0.26*	−0.83*	−0.50*	−0.24*	−0.85*	−0.68*	−1.26*	−1.15*
6	−2.00*	−1.69*	−1.53*	−1.86*	−1.42*	−1.31*	−1.19*	−1.15*
8	0.50*	0.41*	0.05	0.33*	−0.26*	−0.17	−0.49*	−0.75*
10	−0.04	−0.04	−0.30*	−0.09	0.28*	0.06	0.37*	0.30*
12	−0.08	0.06	0.07	0.17	−0.49*	−0.12	−0.67*	−0.60*
14	0.40*	0.55*	−0.10	−0.49*	0.37*	−0.10	0.13	0.41*
16	1.29*	0.87*	0.51*	0.74*	−0.25*	0.04	−0.01	−0.33*
18	0.40*	0.72*	0.27*	0.42*	−0.44*	−0.31*	0.12	0.21*
20	−1.13*	−0.61*	−0.79*	−1.15*	−0.42*	−0.33*	0.08	0.29*
22	0.13	−0.21*	−0.07	0.01	0.53*	−0.11	−0.24*	−0.06
24	−0.99*	−0.53*	−1.28*	−1.32*	−0.03*	−0.08	−0.21*	−0.13*
26	−0.11	−0.16*	−0.24*	0.05	−0.16	−0.25*	−0.26*	−0.37*
28	−0.27*	−0.29*	−0.46*	−0.40*	−0.80*	−0.47*	−0.53*	−0.71*
30	−0.07	0.42*	0.29*	0.12	−0.10*	−0.47*	−0.70*	−0.69*

a big difference in 2008 data as the signs of average returns during 2008 clearly indicate the lack of a systematic pattern. Positive returns are only generated across all information sets for k equal to 0, 16 or 18 basis points on March 31 and only for $k=10$ basis points on April 1. Most of the average out-of-sample returns are statistically significant as indicated by the t -test but the systematic pattern of positive returns from the 2003 data has vanished.

Similar results are obtained for the trading strategy based on the “majority” rule (see Table 7). There are significant and positive returns for k between 0 and 10 basis points for the 2003 samples. Most of the information sets exhibit superior performance to the random walk according to the Anatolyev-Gerko test. Returns across the 2008 samples are, however, not systematically positive and change sign from one information set to another.

There is therefore a pronounced difference between the 2003 and 2008 data sets and also in the performance of the linear model strategy and the genetic algorithm-based strategy when transaction costs are taken into account. This can be explained by the fact that the genetic algorithm chooses the best in-sample strategy according to its returns taking into account transaction costs. In this way the genetic algorithm-based strategy automatically adjusts its trading frequency. In other words, trading rules that tend to trade too often cannot survive in-sample. The linear model strategy cannot adjust to the trading frequency in-sample and therefore it does poorly for small k but improves performance for larger k .

4.1.1. Endogenous inertia parameter

The results reported above show that the profitability of the trading rules critically depends on the value of the inertia parameter. Moreover, the optimal value of k does not stay the same over different sample periods and across different information sets. Therefore it is important to verify profitability of the trading rules based on an ex ante and systematic method for the selection of the inertia parameter. In order to do this we endogenize k in the following way. For each value of the inertia parameter, the genetic algorithm searches for the best in-sample trading rule. In-sample returns are compared with each other, while the rule with the highest return and its corresponding value of k are used to trade out-of-sample. Note, this procedure is using information known to the trader at the time of decision making.¹² The results are provided in Table 8.

Average out-of-sample returns remain positive across all information sets for the first two 2003 samples. During March 17–21, 2003, trading strategies based only on *Screen* and limit order book information generate positive and statistically significant returns while the other two information sets fail to generate positive profit. The samples from 2008 do not produce positive returns for any of the information sets using the endogenously determined value of k .

Table 8 presents the values of the Omega measure of performance (see Shadwick and Keating, 2002) along with Sharpe ratios. Omega is a risk-adjusted performance measure in the sense that it is a ratio of probability weighted gains to losses about a pre-specified threshold, which we take to be zero:

$$\Omega_{\tau} = \left(\int_{\tau}^{\infty} (1 - F(x)) dx \right) / \left(\int_{-\infty}^{\tau} F(x) dx \right),$$

¹²Balvers and Wu (2010) use a dynamic programming framework to design an ex ante optimal filter that maximizes expected returns net of transaction costs.

Table 7

Performance of the GA “majority” trading rule including transaction costs.

This table presents the out-of-sample cumulative returns generated by the “majority” rule based on genetic algorithm for the five different sample periods. 99 independent runs of the genetic algorithm have been performed to select the best in-sample trading rules. The “majority” rule produces a “buy” (sell) signal if the majority of the 99 best in-sample rules produce a “buy” (sell) signal. This combined rule is then used to trade out-of-sample. Returns are adjusted to a daily basis and expressed in percent. k is the threshold value for the trading band—a trader re-evaluates his position only if the current mid-quote is outside the k -band around the last traded price. k is measured in basis points. Transaction costs are reflected in bid–ask spread as trading is based on best bid and ask limit orders. Columns contain results for the different information sets. “Screen” denotes the Screen information, “Book” denotes the limit order book information set, “Order” denotes the order flow set and “Full” corresponds to the combined information set. * indicates significance at the 5% level according to Anatolyev–Gerko test statistics.

k	January 13–17, 2003				February 10–14, 2003				March 17–21, 2003			
	Screen	Book	Order	Full	Screen	Book	Order	Full	Screen	Book	Order	Full
0	1.98*	1.93*	2.23*	2.32*	1.29*	1.30*	0.83*	0.88*	0.67*	0.59*	0.43*	0.54*
2	1.73*	1.79*	1.74*	1.71*	1.14*	1.26*	1.26*	1.13*	0.56*	0.59*	0.51*	−0.09
4	1.60*	1.66*	1.61*	1.64*	1.34*	1.39*	1.29*	1.30*	0.65*	0.57*	0.57*	0.54*
6	1.64*	0.63	0.54	−0.05	1.28*	1.39*	0.99*	1.19*	0.52*	0.54*	0.52*	0.36*
8	1.73*	1.35*	1.10*	1.33*	1.06*	0.93*	0.77*	0.30	0.36*	0.65*	0.04*	0.33*
10	0.69*	0.41	0.87*	0.91*	0.89*	1.08*	0.80*	1.15*	0.24*	0.45*	0.16*	−0.23*
12	−0.16	0.32	0.18	0.75*	0.81*	1.84*	0.88*	0.65*	0.39*	0.21	−0.08	0.16
14	0.49	0.37	0.00	0.77*	1.12*	1.53*	1.34*	1.92*	0.33	0.15*	0.22	0.10
16	−0.16	0.92*	−0.49	−0.17	0.38	1.09*	−0.64*	0.75*	0.19	0.13*	−0.12	−0.59*
18	0.05	0.22	−0.03	0.02	0.53*	0.55*	−0.36*	−0.16	0.14	−0.04	0.07	0.20*
20	−0.32	−0.28	−0.06	0.61	0.44*	0.74*	0.34	−0.06	−0.32	−0.19*	−0.08	−0.11
22	−0.01	−0.39	−0.15	−0.29	−0.03	−0.02	0.09	0.04	0.04	0.13	0.29	0.11
24	0.02	−0.05	0.07	0.14	0.12	−0.55*	0.12	−0.35	0.55*	0.11	0.03	−0.15
26	−0.19	0.16	−0.19*	−0.19*	0.28*	−0.10	−0.41*	−0.43	−0.34	−0.31	−0.04*	−0.24
28	0.08	−0.19	−0.01*	−0.15*	0.22	−0.04	−0.01	0.04	0.06*	−0.22	−0.22	−0.22
30	−0.12	0.17	0.07	−0.02	−0.11	0.01	−0.01	0.20	−0.23	−0.10	−0.21	−0.08

k	March 31, 2008				April 1, 2008			
	Screen	Book	Order	Full	Screen	Book	Order	Full
0	0.13	0.13	0.13	0.13	0.23	−0.72	−0.72	−0.72
2	0.07	0.07	0.07	0.07	−0.72	−0.72	−1.87	−2.01
4	0.88*	−1.42	−2.10*	0.37	−0.82*	−0.72	−1.08*	−1.08
6	−2.25*	−2.43*	−2.49*	−2.43*	−1.81*	−1.53*	−1.39*	−1.39*
8	0.51	0.58	1.09	0.10	−0.96	0.17	−0.46	−1.40
10	0.02	0.45	0.18	−0.04	0.53	−0.64	1.10	0.90
12	0.27	−0.12	−0.61	−0.33	−0.48	0.08	−1.46	−1.08
14	0.82	1.01	0.42	−1.58	1.45	−0.48	0.09	1.20
16	2.57*	2.57*	0.25	0.59	−0.04*	−0.01*	0.35	−0.09
18	0.80	1.65	1.30	0.76	0.45	−0.72	1.14	1.14
20	−1.58	−1.18	−2.00*	−2.44*	−1.14	−0.72	0.39*	0.39*
22	0.60	0.29	−1.67*	0.13	0.43	−1.02	−0.70*	−1.20
24	−1.19	−0.81	−1.45	−1.45	0.17	0.17	0.17	0.17
26	−0.29	−0.29	0.00	−0.29	0.33	−0.76	−0.20	−0.38
28	−0.40	−0.40	−0.40	−0.40	−1.30	−0.44	−0.44	−0.72
30	0.02	0.71	0.08	0.08	0.03	−0.72	−0.72	−0.72

Table 8

Endogenous k , genetic algorithm performance.

This table presents the out-of-sample cumulative returns generated by the genetic algorithm for the five difference sample periods when k is optimally determined based on the in-sample performance. k reflects the inertia threshold value, i.e., the trader re-evaluates his position only if the current mid-quote is outside the k -band around the last traded price. The in-sample period is used to select the best performing trading rule for each k . The best performing rule at k generating the highest in-sample return is used to trade out-of-sample. The exercise is repeated 100 times to generate the empirical distribution of cumulative returns. Returns are adjusted to a daily basis. Transaction costs are reflected in bid–ask spread. Average GA returns column provides average out-of-sample returns, Ω contains the Omega measure, the last panel presents Sharpe ratios of the trading rules. “Screen” denotes the Screen information, “Book” denotes the limit order book information set, “Order” denotes the order flow set and “Full” corresponds to the combined information set. * indicates significance at the 5% level according to t -test statistics.

Sample	Average GA returns				Ω				Sharpe ratio			
	Screen	Book	Order	Full	Screen	Book	Order	Full	Screen	Book	Order	Full
January 13–17, 2003	0.65*	0.57*	0.57*	0.44*	34.9	23.9	15.9	12.9	1.67	1.35	1.31	1.19
February 10–14, 2003	0.98*	0.94*	0.66*	0.66*	96.5	80.2	28.5	53.2	2.81	1.99	1.31	1.28
March 17–21, 2003	0.15*	0.10*	−0.05	−0.04*	2.48	1.73	0.73	0.74	0.39	0.22	−0.11	−0.12
March 31, 2008	−0.65*	−0.09*	−0.51*	−0.68*	0.21	0.82	0.35	0.21	−0.59	−0.08	−0.42	−0.20
April 1, 2008	−0.06	−0.07	−0.15*	−0.18	0.81	0.80	0.56	0.60	−0.08	−0.09	−0.22	−0.63

where F is the cumulative distribution function of returns. As such it reflects the shape of the entire return distribution and all higher moments. The table shows that both the Sharpe and Omega ratios confirm our conclusions from the cumulative returns that there is a substantial difference for all information sets between the 2003 and 2008 data. According to the Omega measure, it is clear that a general bias towards positive returns in 2003 has been replaced, in 2008, by a greater weight being found for negative returns. This is also reflected in the Sharpe ratios.

We now draw together our results from examining Hypothesis 1. First an interesting and important result is the systematic profitability of trading strategies in the FX market for the 2003 data; both the linear model and the genetic algorithm-based approaches generate positive and significant out-of-sample returns even after taking transaction costs into account. These returns could be regarded as a compensation for the risk that traders are exposed to when adopting these strategies. The predictable components in the returns could reflect time-variation in risk premia and the degree of predictability consistent with an efficient market. However, standard performance measures may not reflect the full risks associated with the trading rules. Traders could for instance not only be exposed to market risk as measured by the standard deviation of returns but also to different sources of operational risk.

Another explanation for high out-of-sample returns is simply market inefficiency. The Adaptive Market Hypothesis (AMH) of Lo (2004) states that traders make decisions based on their past experience and learn by receiving negative or positive feedback from the outcomes of those decisions. As a consequence, profitable investment strategies may stop generating positive excess returns because they become more widely exploited using new quantitative methods. Neely, Weller, and Dittmar (1997) found substantial profitability in exchange rate markets during 1974–1995 using a genetic algorithm-based strategy at a

daily frequency. However, in a later study (see Neely and Weller, 2003), they could not confirm this result using 1996 half-hourly data. We are effectively confirming this line of argument on a tick level and with more recent data. We have found that returns drop substantially in the 2008 samples as compared to 2003. We still see positive returns for some combination of inertia parameter and information sets, but when k is endogenized in the trading strategy, we systematically obtain negative performance. The profitability found in the 2003 data is virtually eliminated in the 2008 data. Chaboud, Chiquoine, Hjalmarsson, and Vega (2009) report that during 2003 there was almost no algorithmic high-frequency trading in the FX market, while the fraction of trading volume where at least one of the two counterparties was an algorithmic trader grew up to 60% by the end of 2007. Thus, the profitability of trading exchange rates at the highest frequency appears to have decreased substantially. This observation would be completely consistent with the AMH.

4.2. Hypothesis 2: relative performance

We now turn to Hypothesis 2 and look at the relative performance of the four information sets. In order to test the superior forecasting ability of different conditioning information, we compute the t -statistics of differences between the means of the cumulative returns. This gives us an indication of whether the differential information in each set adds value to the predictions made on the basis of the most basic *Screen* information set. Our main conclusion is that we cannot systematically reject the null hypothesis of the superiority of the *Screen* information. In other words, the enhanced information sets do not appear to add significant value.

The results of the tests for the linear model are provided in Table 9. This table reports values of the Giacomini-White test of conditional superior predictive ability of each of the extended information sets versus the *Screen* information.

In most cases we cannot detect any superior predictive ability for any of the information sets. For most values of the inertia parameter, the extended information sets do not add significant value out-of-sample.

For the genetic algorithm strategies, we compute the distribution of out-of-sample returns and employ the paired t -test to compare the performance through average returns. Table 10 contains the values of the t -statistics.

In Table 10, low values of the inertia parameter *Screen* information appears to be superior to limit order book and the order flow information. There is a k range from 18 to 22 basis points, where the limit order book information dominates the *Screen* information. This mostly appears, however, when the latter generates negative average returns (with the exception of the February 10–14, 2003 sample period).

In addition to the t -test results, we also test for conditional superior predictive ability from the information sets for the genetic algorithm-based strategies. Table 11 provides the values of the Giacomini-White test statistics for the “majority” rule.

In Table 11, there are very few F -statistic, that are statistically significant. There is no clear systematic pattern among those that generate a statistically significant difference in performance between the information sets. In most cases the three information sets based on limit order book variables are not able to significantly outperform the *Screen* information set.

Table 9

Relative performance of the linear model.

The table presents values of Giancomini-White test statistics for comparing the relative out-of-sample performance of the linear model based on three information sets against the *Screen* information. The linear regression coefficients are estimated using in-sample and return predictions are formed for the out-of-sample period. A simple binary trading rule is implemented based on the return predictions. Transaction costs are reflected in the bid–ask spread as trading is based on best bid and ask limit orders. k is the threshold value for the trading band and measured in basis points. Columns “Book-Screen” contains statistics for relative performance of the limit order book versus the *Screen* information, “Order-Screen” compares the order flow versus the *Screen* information and “Full-Screen” corresponds to the combined information set versus the *Screen* information. Asterisks indicate significant values at 5% level; * corresponds to the cases where the *Screen* information outperforms the corresponding information set, ** presents the opposite situation.

k	January 13–17, 2003			February 10–14, 2003			March 17–21, 2003		
	Book-Screen	Order-Screen	Full-Screen	Book-Screen	Order-Screen	Full-Screen	Book-Screen Screen	Order-Screen Screen	Full-Screen Screen
0	5.79	2.20	7.29**	16.70**	1.71	22.02*	6.72**	41.71*	2.46
2	2.17	4.52	1.58	1.46	3.33	1.57	1.46	5.62	6.17*
4	6.04*	1.54	4.77	1.01	3.42	0.39	0.59	1.46	1.78
6	2.03	0.13	0.45	6.38*	0.38	4.13	0.50	1.89	1.89
8	1.19	3.39	1.92	2.01	2.66	0.51	0.69	0.24	0.74
10	N/A	N/A	1.06	0.92	2.75	0.59	0.15	0.13	0.44
12	N/A	1.95	1.46	0.40	N/A	0.40	0.08	N/A	0.13
14	N/A	N/A	N/A	1.59	N/A	1.59	2.34	2.39	3.95
16	N/A	0.82	N/A	0.49	N/A	0.49	N/A	1.09	2.38
18	N/A	2.71	0.82	N/A	2.34	2.34	N/A	2.97	2.97
20	N/A	N/A	2.71	0.77	N/A	1.71	2.47	2.02	4.82
22	N/A	N/A	N/A	N/A	N/A	N/A	2.23	N/A	2.23
24	N/A	N/A	N/A	N/A	2.89	3.43	N/A	N/A	2.11
26	N/A	N/A	N/A	1.95	N/A	1.72	N/A	2.00	2.00
28	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
30	N/A	N/A	N/A	2.99	N/A	1.15	1.07	N/A	1.07

k	March 31, 2008			April 1, 2008		
	Book-Screen	Order-Screen	Full-Screen	Book-Screen	Order-Screen	Full-Screen
0	5.79	2.20	7.29**	10.14*	47.01**	19.23**
2	2.17	4.52	1.58	0.02	4.09	6.95**
4	6.04*	1.54	4.77	2.92	3.39	0.47
6	2.03	0.13	0.45	N/A	2.62	1.78
8	1.19	3.39	1.92	N/A	1.52	2.55
10	N/A	N/A	1.06	N/A	0.04	0.08
12	N/A	1.95	1.46	N/A	5.34	2.49
14	N/A	N/A	N/A	N/A	3.75	3.75
16	N/A	0.82	N/A	N/A	0.77	0.77
18	N/A	2.71	0.82	N/A	0.12	0.12
20	N/A	N/A	2.71	N/A	0.06	1.41
22	N/A	N/A	N/A	N/A	0.79	3.07
24	N/A	N/A	N/A	N/A	1.83	1.83
26	N/A	N/A	N/A	N/A	0.14	0.14
28	N/A	N/A	N/A	N/A	2.00	N/A
30	N/A	N/A	N/A	N/A	2.00	2.00

Table 10

Relative performance of the genetic algorithm model.

The table presents values of paired *t*-test statistics for comparing the relative out-of-sample performance of the genetic algorithm model based on three information sets against the *Screen* information. The in-sample period (the first half of the sample) is then used to select the best performing trading rule and this rule is used to trade out-of-sample (the second half of the sample). The exercise is repeated 100 times to generate the empirical distribution of cumulative returns. Transaction costs are reflected in the bid–ask spread as trading is based on the best bid and ask limit orders. *k* is the threshold value for the trading band and measured in basis points. Columns “Book-Screen” contains statistics for relative performance of the limit order book versus the *Screen* information, “Order-Screen” compares the order flow versus the *Screen* information and “Full-Screen” corresponds to the combined information set versus the *Screen* information. Asterisks indicate significant values at 5% level; * corresponds to the cases where the *Screen* information outperforms the corresponding information set, ** presents the opposite situation.

<i>k</i>	January 13–17, 2003			February 10–14, 2003			March 17–21, 2003		
	Book-Screen	Order-Screen	Full-Screen	Book-Screen	Order-Screen	Full-Screen	Book-Screen Screen	Order-Screen Screen	Full-Screen Screen
0	−0.46	−1.17	−1.54	0.12	−2.72*	−2.09*	−3.84*	−6.99*	−7.61*
2	−0.22	−2.37*	−1.85*	−0.20	−0.55	−1.59	−3.49*	−5.58*	−9.88*
4	−2.73*	−3.34*	−2.95*	0.22	−2.82*	−3.04*	−2.20*	−5.21*	−6.58*
6	−3.05*	−4.51*	−5.88*	−2.52*	−3.98*	−7.71*	−2.26*	−4.68*	−5.92*
8	−1.43	−1.85*	−2.57*	−2.31*	−5.01*	−6.69*	−0.09	−4.43*	−4.12*
10	−1.88*	−0.03	−1.79*	−0.77	−2.98*	−0.61	0.44	−1.75*	−2.13*
12	2.09**	−0.11	1.93**	−0.15	−5.35*	−3.54*	−1.26	−1.66	−2.35*
14	3.01**	−2.20*	0.51	0.26	−2.69*	−0.45	−0.86	0.14	−0.28
16	5.18**	−1.58	1.12	0.66	−8.66*	−2.30*	−0.47	−2.16*	−3.85*
18	2.55**	−1.27	−0.32	−0.72	−8.12*	−4.94*	−2.70*	1.02	−0.31
20	0.11	3.09**	3.78**	2.36**	2.00**	1.49	−0.41	3.87**	2.68**
22	−1.97*	−0.66	−1.57	3.37**	1.19	1.92**	0.69	−0.23	0.74
24	−0.73	1.73**	0.43	−7.24*	0.31	−2.50*	−6.34*	−1.94*	−3.97*
26	0.84	2.07**	2.16**	−2.69*	−6.71*	−5.30*	−2.38*	−1.06	−0.82
28	−3.00*	0.38	−1.36	−1.77*	0.58	−0.89	0.23	−0.32	−0.80
30	0.24	0.45	0.26	−0.64	0.32	−1.14	−3.42*	1.27	−1.00

<i>k</i>	March 31, 2008			April 1, 2008		
	Book-Screen	Order-Screen	Full-Screen	Book-Screen	Order-Screen	Full-Screen
0	−1.96*	−2.63*	−3.58*	−2.09*	−6.09*	−4.20*
2	−1.56	−1.82*	−1.70*	0.51	−22.85*	−22.67*
4	−5.96*	−4.50*	−3.26*	1.71**	−4.88*	−3.84*
6	2.18**	2.54*	0.99	0.84	1.58	2.66**
8	−0.68	−2.30*	−1.07	0.60	−1.63	−3.65*
10	−0.02	−1.75*	−0.34	−1.39	0.52	0.11
12	0.78	0.88	1.46	2.52**	−1.04	−0.76
14	0.98	−2.80*	−5.26*	−3.03*	−1.62	0.27
16	−1.97*	−4.11*	−2.79*	2.98**	2.43**	−0.86
18	2.33**	−0.82	0.11	1.33	5.22**	5.77**
20	3.32**	2.30**	−0.14	0.87	5.39**	6.94**
22	−2.11*	−1.28	−0.76	−3.65*	−4.41*	−3.31*
24	2.71**	−1.84*	−2.30*	−0.45	−1.70*	−1.02
26	−0.42	−1.08	1.27	−0.82	−0.80	−1.89*
28	−0.15	−2.01*	−1.20	6.17**	4.64**	1.33
30	3.62**	2.70**	1.37	−4.70*	−6.87*	−6.93*

Table 11

Relative performance of the GA “majority” trading rule.

The table presents values of Giancomini-White test statistics for comparing the relative out-of-sample performance of the “majority” rule based on three information sets against the *Screen* information. 99 independent runs of the genetic algorithm have been performed to select the best in-sample trading rules. The “majority” rule produces a “buy” (sell) signal if the majority of the 99 best in-sample rules produce a “buy” (sell) signal. This combined rule is then used to trade out-of-sample. Transaction costs are reflected in the bid–ask spread as trading is based on the best bid and ask limit orders. k is the threshold value for the trading band and measured in basis points. Columns “Book-Screen” contains statistics for relative performance of the limit order book versus the *Screen* information, “Order-Screen” compares the order flow versus the *Screen* information and “Full-Screen” corresponds to the combined information set versus the *Screen* information. Asterisks indicate significant values at 5% level; * corresponds to the cases where the *Screen* information outperforms the corresponding information set, ** presents the opposite situation.

k	January 13–17, 2003			February 10–14, 2003			March 17–21, 2003		
	Book-Screen	Order-Screen	Full-Screen	Book-Screen	Order-Screen	Full-Screen	Book-Screen Screen	Order-Screen Screen	Full-Screen Screen
0	3.09	1.78	1.95	0.00	2.16	2.11	0.27	2.74	2.11
2	1.13	0.06	0.19	0.56	0.42	1.12	2.32	0.34	3.91
4	1.60	0.00	0.79	0.51	0.81	0.24	0.68	0.68	1.52
6	5.96	8.52*	10.13*	2.19	8.76*	3.31	1.48	1.40	1.90
8	2.43	3.37	2.03	0.23	0.55	4.44	3.90	0.32	1.93
10	1.09	0.33	0.36	1.63	0.13	1.46	1.63	0.03	1.67
12	1.74	0.70	3.24	8.16**	0.45	0.44	0.36	1.23	3.36
14	0.78	2.89	2.18	3.94	1.23	7.42**	2.73	0.07	1.24
16	5.18	3.47	3.07	3.05	8.52	1.27	0.20	2.94	3.06
18	1.41	0.82	0.52	0.26	6.96	5.06	0.38	0.22	0.22
20	0.90	1.38	2.35	6.05**	1.80	7.26**	0.70	2.58	2.95
22	2.41	2.86	1.42	3.14	1.61	2.05	1.51	0.59	0.41
24	0.33	1.67	0.01	3.43	1.87	4.09	1.17	2.72	3.93
26	5.05	0.01	0.01	1.18	3.70	3.81	0.59	0.86	2.00
28	1.33	0.52	1.66	1.27	3.37	2.76	1.31	1.31	1.31
30	0.69	3.13	4.02	2.74	2.55	1.09	2.00	N/A	1.07

k	March 31, 2008			April 1, 2008		
	Book-Screen	Order-Screen	Full-Screen	Book-Screen	Order-Screen	Full-Screen
0	1.08	1.18	1.13	0.99	0.99	0.99
2	1.35	1.30	1.65	1.26	10.70*	11.34*
4	3.54	5.51	0.84	2.00	0.94	0.94
6	1.07	1.87	1.07	0.08	1.55	1.34
8	0.27	2.25	1.77	3.18	3.69	1.72
10	0.46	1.08	1.18	3.05	2.23	1.05
12	0.26	0.36	0.01	1.00	5.10	5.74
14	1.16	1.11	4.99	3.69	1.97	2.49
16	N/A	3.70	3.04	N/A	2.00	0.96
18	2.86	0.07	0.85	0.11	2.14	2.14
20	0.06	1.01	1.72	2.00	1.96	1.96
22	2.70	3.67	1.32	2.76	2.25	2.97
24	1.75	0.64	1.78	N/A	N/A	N/A
26	2.44	N/A	N/A	2.00	0.04	2.31
28	N/A	N/A	N/A	2.00	2.00	2.00
30	2.00	2.30	2.30	N/A	N/A	N/A

These results again suggest that the limit order book information does not appear to carry significant additional information over that included in past price and quantities, that can be systematically profitable out-of-sample. This result would be in line with a large theoretical literature claiming that informed traders use market orders to exploit their private information. This means that limit orders and the structure of the book does not carry any substantial information that can be exploited out-of-sample. Although we confirm that order book information has statistical in-sample explanatory power, this predictability cannot be systematically transferred into economically significant profit, at least beyond that which is already in the past prices.

5. Conclusions

In this paper we examine the predictability and profitability of the U.S. dollar sterling exchange rate using limit order book information. We test formally the hypothesis of whether the limit order book information can be profitably exploited out-of-sample for five different samples during 2003 and 2008. Two approaches are used to construct trading strategies: linear regression and a genetic algorithm.

We show that there is a high level of profitability during 2003 and trading strategies generate positive and significant out-of-sample returns net of transaction costs. The level of profitability appears to drop considerably if not being eliminated completely in the more recent 2008 data. This finding is in line with the Adaptive Market Hypothesis given that there has been a dramatic rise in algorithmic trading activity since 2003. To the best of our knowledge, this gain in efficiency has not yet been reported in the literature for the FX markets but would be consistent with the analysis by [Hendershott, Jones, and Menkveld \(2010\)](#) on the impact of algorithmic trading for equity markets.

We do not find any systematic evidence that limit order book information can add significant economic value to the out-of-sample performance of the trading strategies. We look at four information sets, the first of which is based purely on past price information and the quantities of best limit orders which is visible on the screens of trading platforms. The other three are based on the limit order book and order flow information in addition to the screen information set. The information contained in the three enhanced information sets seem to be not robust enough to significantly contribute to the profitability of the trading strategies.

Our results suggest that the advent of algorithmic trading has had a huge effect on the efficiency of financial markets. The theoretical models reviewed in [Section 2](#) have been developed when humans were taking trading decisions in FX markets. [Chaboud, Chiquoine, Hjalmarsson, and Vega \(2009\)](#) argue that strategies generated by computers are much more correlated among themselves than humans decisions are. This clearly shows the need to develop new models as to how the algorithmic trading rules are being designed and the extent to which limit orders and the structure of the order book are being exploited in these trading schemes.

References

- Alexander, S., 1961. Price movements in speculative markets: trends or random walks. *Industrial Management Review* 2, 7–26.
- Allen, F., Karjalainen, R., 1999. Using genetic algorithms to find technical trading rules. *Journal of Financial Economics* 51, 245–271.

- Anatolyev, S., Gerko, A., 2005. A trading approach to testing for predictability. *Journal of Business & Economic Statistics* 23, 455–462.
- Arifovic, J., 1996. The behavior of the exchange rate in the genetic algorithm and experimental economies. *Journal of Political Economy* 104, 510–541.
- Balvers, R., Wu, Y., 2010. Optimal transaction filters under transitory trading opportunities: theory and empirical illustration. *Journal of Financial Markets* 13, 129–156.
- Bloomfield, R., O'Hara, M., Saar, G., 2005. The “make or take” decision in an electronic market: evidence on the evolution of liquidity. *Journal of Financial Economics* 75, 165–199.
- Cao, C., Hansch, O., Wang, X., 2009. The information content of an open limit order book. *The Journal of Futures Markets* 29, 16–41.
- Chaboud, A., Chiquoine, B., Hjalmarsson, E., Vega, C., 2009. Rise of the machines: algorithmic trading in the foreign exchange market. Working Paper, Federal Reserve System.
- Chakravarty, S., Holden, C., 1995. An integrated model of market and limit orders. *Journal of Financial Intermediation* 4, 213–241.
- Chen, S., Duffy, J., Yeh, C., 1999. Genetic programming in the coordination game with a chaotic best-response function, in: *Proceedings of the 1996 Evolutionary Programming Conference*, San Diego, CA.
- Chen, S., Yeh, C., 1997a. On the coordination and adaptability of the large economy: an application of genetic programming to the cobweb model. In: Angeline, P., Kinnear, K. (Eds.), *Advances in Genetic Programming II*. MIT Press, Cambridge MA.
- Chen, S., Yeh, C., 1997b. Toward a computable approach to the efficient market hypothesis: an application of genetic programming. *Journal of Economic Dynamics and Control* 21, 1043–1063.
- Chordia, T., Subrahmanyam, A., 2004. Order imbalance and individual stock returns: theory and evidence. *Journal of Financial Economics*, 485–518.
- Cooper, M., 1999. Filter rules based on price and volume in individual security overreaction. *Review of Financial Studies* 12, 901–935.
- Danielsson, J., Luo, J., Payne, R., 2002. Exchange rate determination and inter-market order flow effects. Unpublished manuscript.
- Della Corte, P., Sarno, L., Tsiakas, I., 2009. An economic evaluation of empirical exchange rate models. *Review of Financial Studies* 22, 3491–3530.
- Dempster, M., Jones, C., 2001. A real time adaptive trading system using genetic programming. *Quantitative Finance* 1, 397–413.
- Dworman, G., Kimbrough, S., Laing, J., 1996. On automated discovery of models using genetic programming: bargaining in a three-agent coalitions game. *Journal of Management Information Systems* 12, 97–125.
- Easley, D., O'Hara, M., 2010. Liquidity and valuation in an uncertain world. *Journal of Financial Economics* 97, 1–11.
- Evans, M., Lyons, R., 2005. Meese and Rogoff redux: micro-based exchange rate forecasting. *American Economic Review Papers and Proceedings* 95, 405–414.
- Evans, M., Lyons, R., 2006. Understanding order flow. *International Journal of Finance and Economics* 11, 3–23.
- Fama, E., Blume, M., 1966. Filter rules and stock market tradings. *Journal of Business* 39, 226–241.
- Foucault, T., 1999. Order flow composition and trading costs in a dynamic limit order market. *Journal of Financial Markets* 2, 99–134.
- Froot, K., Ramadorai, T., 2005. Currency returns, intrinsic value and institutional investor flows. *Journal of Finance* 60, 1535–1566.
- Giacomini, R., White, H., 2006. Tests of conditional predictive ability. *Econometrica* 74, 1545–1578.
- Glosten, L., 1994. Is the electronic open limit order book inevitable? *Journal of Finance* 49, 1127–1161.
- Granger, C., Machina, M., 2006. Forecasting and decision theory. In: Elliott, G., Granger, C., Timmermann, A. (Eds.), *Handbook of Economic Forecasting*. Elsevier.
- Granger, C., Pesaran, H., 2000. Economic and statistical measures of forecast accuracy. *Journal of Forecasting* 19, 537–560.
- Handa, P., Schwartz, R., 1996. Limit order trading. *Journal of Finance* 51, 1835–1861.
- Handa, P., Schwartz, R., 2003. Quote setting and price formation in an order driven market. *Journal of Financial Markets* 6, 461–489.
- Hansen, P., 2010. In-sample fit and out-of-sample fit: their joint distribution and its implications for model selection. Working Paper, Stanford University.
- Harris, L., 1998. Optimal dynamic order submission strategies in some stylized trading problems. *Financial Markets, Institutions & Instruments* 7, 1–76.

- Harris, L., Hasbrouck, J., 1996. Market vs. limit orders: the SuperDOT evidence on order submission strategy. *Journal of Financial and Quantitative Analysis* 31, 213–231.
- Harris, L., Panchapagesan, V., 2005. The information content of the limit order book: evidence from the NYSE specialist trading decisions. *Journal of Financial Markets* 8, 25–67.
- Hellström, J., Simonsen, O., 2009. Does the open limit order book reveal information about short-run stock price movements? Unpublished manuscript.
- Hendershott, T., Jones, C., Menkveld, A., 2010. Does algorithmic trading improve liquidity? *Journal of Finance* 66, 1–33.
- Hillman, R., Salmon, M., 2007. Intrinsic stationarity: investigating predictability in real-time FOREX transactions. *Journal of Financial Forecasting* 1, 3–43.
- Huang, R., Stoll, H., 1994. Market microstructure and stock return predictions. *Review of Financial Studies* 7, 179–213.
- Kaniel, R., Liu, H., 2006. What orders do informed traders use? *Journal of Business* 79, 1867–1913.
- Kavajecz, K., 1999. A specialist's quoted depth and the limit order book. *Journal of Finance* 54, 747–771.
- Knez, P., Ready, M., 1996. Estimating the profits from trading strategies. *Review of Financial Studies* 9, 1121–1163.
- Latza, T., Payne, R., 2010. Measuring the information content of limit order flows. Unpublished manuscript.
- Leitch, G., Tanner, E., 1991. Economic forecast evaluation: profits versus the conventional error measures. *American Economic Review* 81, 580–590.
- Lo, A., 2004. The adaptive markets hypothesis. *Journal of Portfolio Management* 30, 14–29.
- Neely, C., Weller, P., 2001. Technical analysis and central bank intervention. *Journal of International Money and Finance* 20, 949–970.
- Neely, C., Weller, P., 2003. Intraday technical trading in the foreign exchange market. *Journal of International Money and Finance* 22, 223–237.
- Neely, C., Weller, P., Dittmar, R., 1997. Is technical analysis in the foreign exchange market profitable? *Journal of Financial and Quantitative Analysis* 32, 405–426.
- Nix, A., Vose, M., 1992. Modeling genetic algorithms with Markov chains. *Annals of Mathematics and Artificial Intelligence* 5, 79–88.
- Parlour, C., 1998. Price dynamics in limit order markets. *Review of Financial Studies* 11, 789–816.
- Pesaran, H., Skouras, S., 2002. Decision-based methods for forecast evaluation. *Companion to Economic Forecasting*. Blackwell Publishers, pp. 241–267.
- Ranaldo, A., 2004. Order aggressiveness in limit order book markets. *Journal of Financial Markets* 7, 53–74.
- Rime, D., Sarno, L., Sojli, E., 2010. Exchange rate forecasting, order flow and macroeconomic information. *Journal of International Economics* 80, 72–88.
- Rock, K., 1996. The specialist's order book and price anomalies. Working Paper, Harvard University.
- Rosu, I., 2010. Liquidity and Information in Order Driven Markets. SSRN Working Paper Series.
- Sager, M., Taylor, M., 2008. Commercially available order flow data and exchange rate movements: Caveat emptor. *Journal of Money, Credit and Banking* 40, 583–625.
- Satchell, S., Timmermann, A., 1995. An assessment of the economic value of non-linear foreign exchange rate forecasts. *Journal of Forecasting* 14, 477–497.
- Seppi, D., 1997. Liquidity provision with limit orders and a strategic specialist. *Review of Financial Studies* 10, 103–150.
- Shadwick, W., Keating, C., 2002. A universal performance measure. *Journal of Performance Measurement* 6, 59–85.
- Vose, M., 1993. Modeling simple genetic algorithms. In: Whitley, L. (Ed.), *Foundations of Genetic Algorithms*, vol. 2. Morgan Kaufman.